

MongoDB AI Arena - Application Architecture Overview

Based on my exploration of the codebase, here's how the different components fit together:

High-Level Purpose

This is a cloud-based training platform that helps users learn MongoDB through hands-on exercises. It's themed around building an Airbnb-like short-term rental application where participants complete coding challenges, compete on a leaderboard, and learn MongoDB concepts like CRUD operations, aggregations, search, and vector search.

Component Breakdown

1. Infrastructure Layer (`utils/`)

This contains all the scripts and configuration for spinning up the cloud environment:

`utils/arena-terraform/` - Main Orchestration

- Purpose: Terraform/Terragrunt configuration to deploy the entire infrastructure
- Key Files:
 - `config.yaml` - Configuration per customer (MongoDB credentials, AWS settings, domain)
 - `airbnb/` - Template directory that gets duplicated per customer/event
 - `scenario.json` - Defines which exercises/modules to include for this event
- What it creates: AWS EKS cluster, MongoDB Atlas cluster, networking, SSL certificates

`utils/atlas-cluster/` - MongoDB Setup

- Purpose: Creates MongoDB Atlas cluster and populates it with sample data
- Creates:
 - Atlas cluster with M30+ tier
 - Database users for participants (from `user_list.csv`)
 - Indexes for CRUD, Search, and Vector Search exercises
 - Loads sample Airbnb dataset

utils/eks-cluster/ - Kubernetes Resources

- Purpose: Deploys all Kubernetes services to AWS EKS
 - Deploys:
 - VSCode Online (`mdb-openvscode/`) - Browser-based IDE for participants
 - Arena Portal (`mongodb-arena-portal/`) - Landing page to claim credentials
 - Documentation (`docs-nginx/`) - Hosts the Jekyll docs site
 - LiteLLM (`litellm/`) - AI proxy for Claude/GPT integration
 - Results Processor (`results-processor/`) - Validates exercise submissions
 - Nginx - Reverse proxy and load balancer
 - Redis - Caching for LLM requests
-

2. Documentation Layer (`docs/`)

This is the instructional content participants see:

- Jekyll-based static site deployed to the EKS cluster
 - Structure:
 - `_pages/crud/` - CRUD operation exercises (8 exercises)
 - `_pages/pipeline/` - Aggregation pipeline exercises
 - `_pages/search/` - Atlas Search exercises
 - `_pages/vector-search/` - Vector search with AI chatbot
 - `_pages/vibe-coding/` - AI-assisted coding exercises
 - `_pages/guided/` - Guided workshop instructions
 - Hosted at: `https://{customer}.mongoarena.com` (configured via nginx)
-

3. Sample Application (The "Airbnb Clone")

This is the working application participants modify during exercises:

app/ - Frontend (Next.js)

- Purpose: User-facing Airbnb-like rental search interface
- Tech: Next.js 15, React 19, TailwindCSS
- Features:
 - Property listings with filters (amenities, beds, property type)
 - Map view (Leaflet/React-Leaflet)

- Search functionality
- Host analytics dashboard
- Exercise status indicators (shows which exercises are complete)
- Chatbot interface (vector search exercise)
- Runs on: Port 3000 in participant's VSCode Online environment

server/ - Backend (Node.js/Express)

- Purpose: API server that participants implement exercises in
 - Tech: Express.js, MongoDB Node Driver, LangChain
 - Key Routes:
 - `/api/listingsAndReviews` - Property CRUD operations
 - `/api/results` - Exercise validation and leaderboard
 - `/api/chat` - AI chatbot (vector search)
 - Exercise Files:
 - `server/src/lab/*.lab.js` - Incomplete exercise files participants edit
 - `utils/answers/*.lab.js` - Solution files for validation
 - Testing: Mocha tests validate implementations and submit scores
-

4. Arena Portal (`utils/eks-cluster/mongodb-arena-portal/`)

This is the participant onboarding system:

- Frontend (Next.js/React/TypeScript): Landing page where participants claim their credentials
 - Backend (Flask/Python):
 - Manages participant database
 - Tracks which credentials are "taken"
 - Shows available slots and leaderboard
 - Database: Uses MongoDB to track participants
 - Hosted at: Root domain (e.g., <https://airbnb.mongoarena.com>)
-

5. Testing & Validation System

The automatic exercise grading system:

- `server/src/saveTestResults.js`: Runs Mocha tests and posts results
- Results Processor (Java/Spring Boot in EKS):
 - Receives test results

- Validates against answer key
 - Updates leaderboard
 - Tracks timing for "timed" mode or scores for "score" mode
 - Exercise Status Component (`app/src/components/ExerciseStatus.js`): Shows green/red indicators in the frontend
-

Data Flow

1. SA runs ``terraform apply``
 - Creates AWS EKS + Atlas Cluster
 - Deploys all services (VSCode, Portal, Docs, etc.)
 2. Participant visits `{customer}.mongoarena.com`
 - Claims credentials from Arena Portal
 - Gets username/password for VSCode Online
 3. Participant opens VSCode Online
 - Pre-loaded with app + server code
 - Follows documentation site for exercises
 4. Participant edits `server/src/lab/*.lab.js` files
 - Frontend app at `localhost:3000` shows results
 - Runs `npm test` to validate solution
 5. Test results → Results Processor
 - Validates against answer key
 - Updates leaderboard
 - Shows green checkmark in frontend
-

Directory Purpose Summary

Directory	Purpose	Who Uses It
<code>utils/arena-terragrunt/</code>	Infrastructure setup scripts	MongoDB SAs only
<code>utils/atlas-cluster/</code>	MongoDB cluster config	MongoDB SAs only
<code>utils/eks-cluster/</code>	Kubernetes deployments	MongoDB SAs only
<code>docs/</code>	Exercise instructions (Jekyll)	Participants view this
<code>app/</code>	Frontend sample app	Participants see/use this
<code>server/</code>	Backend API (exercise files)	Participants edit this
<code>utils/answers/</code>	Solution keys	Validation only
<code>utils/eks-cluster/mongodb-arena-portal/</code>	Credential management portal	Participants claim credentials

Key Technologies

- Cloud: AWS EKS (Kubernetes), Route53 (DNS), S3 (Terraform state)
- Database: MongoDB Atlas (M30+)
- IaC: Terraform + Terragrunt
- Container Orchestration: Kubernetes + Helm charts
- Frontend: Next.js 15, React 19, TailwindCSS
- Backend: Node.js/Express, Python/Flask, Java/Spring Boot
- AI/LLM: LiteLLM proxy, LangChain, Claude/GPT

- Documentation: Jekyll (GitHub Pages style)
- IDE: VSCode in browser (OpenVSCode Server)